

TCS-503: Design and Analysis of Algorithms

Unit II

Advanced Data Structures:


B-Tree

Unit II

- Advanced Data Structures:
 - Red-Black Trees
 - Augmenting Data Structure
 - **B-Trees**
 - Binomial Heaps
 - Fibonacci Heaps
 - Data Structure for Disjoint Sets

What is rooted tree?

A rooted tree is a free tree in which one of the vertices is distinguished from the others. The distinguished vertex is called the root of the tree.



B-Tree

A **B-Tree** T is a **rooted tree** having the following properties:

1. Every **node** x has the following **fields**:

- $n[x]$, the number of keys currently stored in node x .
- the $n[x]$ keys themselves, stored in **increasing order**, so that

$$key_1[x] \leq key_2[x] \leq \dots \leq key_{n[x]}[x].$$

- $leaf[x]$, a Boolean value that is **TRUE**: if x is a leaf and **FALSE** if x is an **internal node**.

$n[x]$	$key_1[x]$	$key_2[x]$	$key_3[x]$	$key_4[x]$	$key_5[x]$	$key_6[x]$	$key_7[x]$	$key_8[x]$	$leaf[x]$
$c_1[x]$	$c_2[x]$	$c_3[x]$	$c_4[x]$	$c_5[x]$	$c_6[x]$	$c_7[x]$	$c_8[x]$	$c_9[x]$	

B-Tree

2. Each **internal node** x also contains $n[x] + 1$ pointers $c_1[x], c_2[x], \dots, c_{n[x] + 1}[x]$ to its children. **Leaf nodes** have no children, so their c_i fields are undefined.


3. The keys $key_i[x]$ separates the ranges of keys stored in each sub-tree: if k_i is any key stored in the sub-tree with root $c_i[x]$, then

$$k_1 \leq key_1[x] \leq k_2 \leq key_2[x] \leq \dots \leq key_{n[x]}[x] \leq k_{n[x] + 1}$$

4. **All leaves** have the **same depth**, which is the tree's height h .


m-way Search Tree

$m = 4$




It is not balanced!

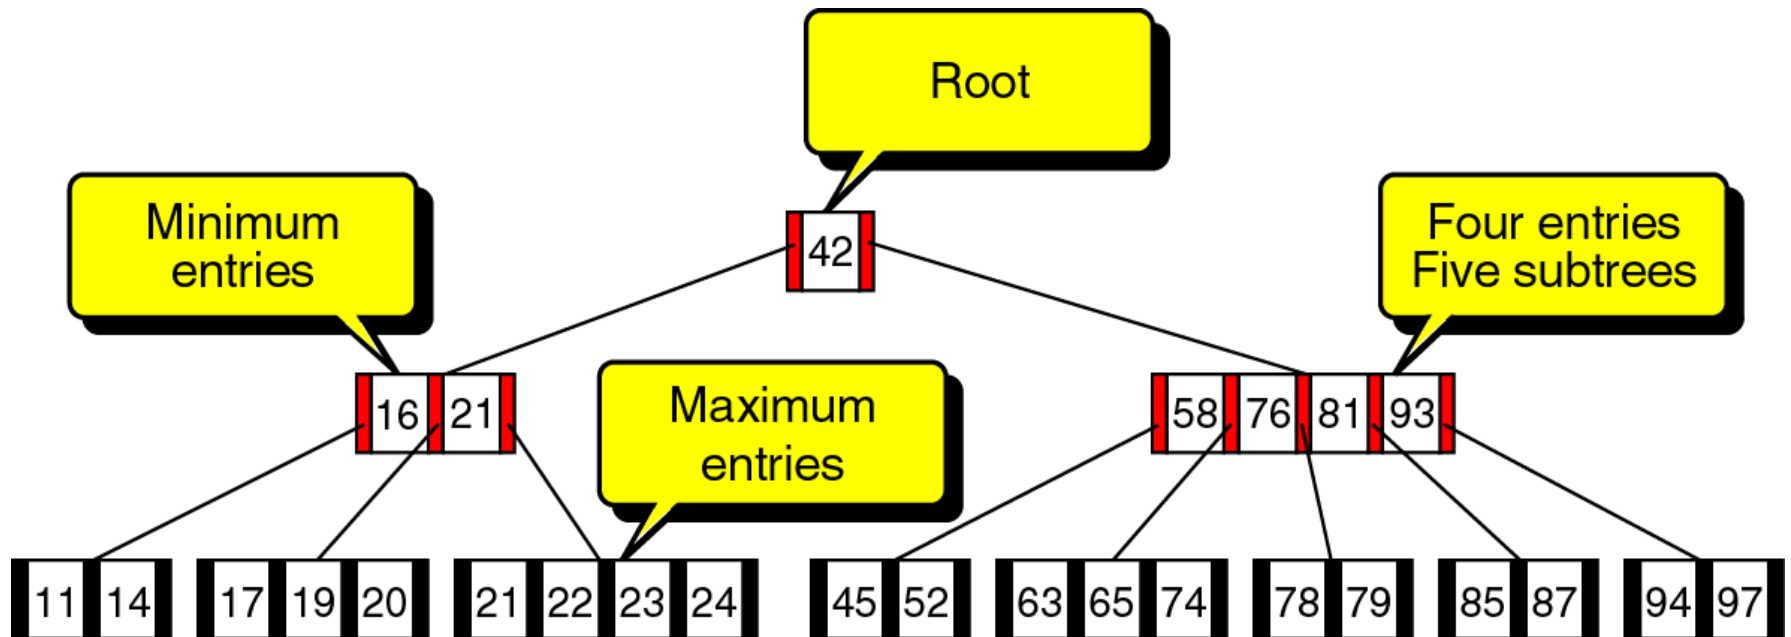
m-way Search Tree



B-Tree: Perfectly balanced (m=) 4-way search tree



B-Tree: Perfectly Balanced m- way search Tree



A B-tree of order; $m=5$.


Min entry : $\text{Ceiling}[5/2] - 1 = 2$ entries.

Max entry: $5 - 1 = 4$ entries.

Min subtrees : $\text{Ceiling}[5/2]=3$

Max subtrees: 5

B-Tree



An internal node x containing $n[x]$ keys has $n[x] + 1$ children.

All leaves are at the same depth in the tree.

B-Tree

5. There are **lower and upper bounds** on the number of keys a node can contain.

These bounds can be expressed in terms of a fixed integer $t \geq 2$ called the **minimum degree of the B-Tree**:

a. Every node other than the root must have at least $t-1$ keys.

Every Internal node other than the root thus has at least t children .


If the tree is **non-empty**, the root must have at least one key.

b. Every node can contain at most $2t - 1$ keys.

Therefore, an internal node can have at most $2t$ children.

We say that a node is full if it contains exactly $2t - 1$ keys.


B-Tree



B-Tree of height $h=3$ containing a minimum possible number of keys.

Shown inside each node x is $n[x]$.

B-Tree



A **B-tree** of height 2 containing over one billion keys. Each internal node and leaf contains 1000 keys. There are 1001 nodes at depth 1 and over one million leaves at depth 2. Shown inside each node x is $n[x]$, the number of keys in x .

B-Tree

The simplest B-Tree occurs when $t=2$ (Binary Tree).

Every internal node then has either 2, 3, or 4 children, and we have a *2-3-4 tree*.

Learn DAA: From B K Sharma

B-Tree

$x \leftarrow$ *a pointer to some object*

Basic Operation on B-Tree

$x \leftarrow$ a pointer to some object

B-TREE-SEARCH(x , k)

B-TREE-CREATE(T)

B-TREE-INSERT(T , k)

B-TREE-DELETE(x , k)

Search Operation on B-Tree


B-TREE-SEARCH(root[T], R)

Searching a B-tree is much like searching a binary search tree, except that instead of making a binary, or "two-way," branching decision at each node, we make a multi-way branching decision according to the number of the node's children.

More precisely, at each internal node x , we make an $(n[x] + 1)$ -way branching decision.

Search Operation on B-Tree

B-TREE-SEARCH(root[T], R)



Insertion Operation on B-Tree

We insert the new key into an existing leaf node.

Since we cannot insert a key into a leaf node that is full, we introduce an operation that *splits a full node y* (having $2t - 1$ keys) around its *median key $key_{\lceil t \rceil}[y]$* into *two nodes having $t - 1$ keys each*.


The median key moves up into y 's parent to identify the dividing point between the two new trees.

But if y 's parent is also full, it must be split before the new key can be inserted, and thus this need to split **full nodes** can propagate all the way up the tree.


Insertion Operation on B-Tree

The minimum degree t for this B-tree is 3 ($t=3$), so a node can hold at most 5 keys.


(a) Initial Tree




(b) Insert B




Insertion Operation on B-Tree




(c) Insert Q




Insertion Operation on B-Tree




(d) Insert L



Insertion Operation on B-Tree



(e) Insert F






Insertion Operation on B-Tree

Insert the following keys into an empty B-Tree of order $t=5$ in order: 1 12 8 2 25 6 14 28 17 7 52 16 48 68 3 26 29 53 55 45

The first four items go into the root:




Insert 25



6, 14, 28 get added to the leaf nodes:

Insertion Operation on B-Tree


1 12 8 2 25 6 14 28 **17** 7 52 16 48 68 3 26 29 53 55
45



Insert 17


Inserting 17 to the right leaf node would over-fill it, so we take the middle key, promote it (to the root) and split the leaf

12, 14, **17**, 25, 28




Insertion Operation on B-Tree

1 12 8 2 25 6 14 28 17 7 52 16 48 68 3 26 29 53 55
45




7, 52, 16, 48 get added to the leaf nodes



Insertion Operation on B-Tree


1 12 8 2 25 6 14 28 17 7 52 16 48 68 3 26 29 53 55
45



Insert 68


Inserting 68 causes us to split the right most leaf, promoting 48 to the root:

25, 28, 48, 52, 68




Insertion Operation on B-Tree

1 12 8 2 25 6 14 28 17 7 52 16 48 68 3 26 29 53 55
45




Insert 3

Inserting 3 causes us to split the left most leaf, promoting 3 to the root: [1, 2, 3, 6, 7]




Insertion Operation on B-Tree

1 12 8 2 25 6 14 28 17 7 52 16 48 68 3 26 29 53 55
45




26, 29, 53, 55 then go into the leaves:




Insertion Operation on B-Tree

1 12 8 2 25 6 14 28 17 7 52 16 48 68 3 26 29 53 55
45




Insert 45

Inserting 45 causes a split of




and promoting 28 to the root



Learn DAA: From B K Sharma

Insertion Operation on B-Tree

1 12 8 2 25 6 14 28 17 7 52 16 48 68 3 26 29 53 55
45



Exercise in Inserting a B-Tree

Insert the following keys to a 5-way B-tree:

3, 7, 9, 23, 45, 1, 5, 14, 25, 24, 13, 11, 8, 19, 4,
31, 35, 56

Show the results of inserting the keys

*F, S, Q, K, C, L, H, T, V, W, M, R, N, P, A,
B, X, Y, D, Z, E*

in order into an empty B-tree with minimum degree 2.

Only draw the configurations of the tree just before some node must split, and also draw the final configuration.

Deletion in a B-Tree

B-tree Delete distinguishes three different stages/scenarios for deletion:

Case 1: key k found in leaf node


Case 2: key k found in internal node

Case 3: key k suspected in lower level node

Deletion in a B-Tree


The minimum degree for this B-tree is $t = 3$, so a node (other than the root) cannot have fewer than 2 keys.

Initial Tree




Deleting Keys

Initial Tree



Case 1: If the key k is in node x , and x is a leaf, delete k from x

Delete F



Deleting Keys

Case 2: If the key k is in node x , and x is not a leaf, delete k from x

a) If the child y that precedes k in node x has at least t keys, then find the predecessor k' of k in the sub-tree rooted at y . Recursively delete k' , and replace k with k' in x . (Finding k and deleting it can be performed in a single downward pass.)

b) Symmetrically, if the child z that follows k in node x has at least t keys, then find the successor k of k in the subtree rooted at z . Recursively delete k , and replace k by k in x . (Finding k and deleting it can be performed in a single downward pass.)

Deleting Keys


c) Otherwise, if both y and z have only $t - 1$ keys, merge k and all of z into y , so that x loses both k and the pointer to z , and y now contains $2t - 1$ keys. Then, free z and recursively delete k from y .

Deleting Keys


Case 2a: If the child y that precedes k in node x has at least t keys, then find the predecessor k' of k in the sub-tree rooted at y . Recursively delete k' , and replace k with k' in x .

Delete M

This is case 2a: the predecessor L of M is moved up to take M 's position.




Deleting Keys



Deleting Keys

Delete G .




This is case 2c: if both y and z have only $t - 1$ keys, merge k and all of z into y , so that x loses both k and the pointer to z , and y now contains $2t - 1$ keys. Then, free z and recursively delete k from y .

Deleting Keys


Delete *G*.

G is pushed down to make node *DEGJK*, and then *G* is deleted from this leaf (case 1).



Deleting Keys

Delete G.



Deleting Keys

Case 3. If the key k is not present in internal node x , determine the root $c_i[x]$ of the appropriate subtree that must contain k , if k is in the tree at all. If $c_i[x]$ has only $t - 1$ keys, execute step 3a or 3b as necessary to guarantee that we descend to a node containing at least t keys. Then, finish by recursing on the appropriate child of x .

Deleting Keys

- a. If $c_i[x]$ has only $t - 1$ keys but has an immediate sibling with at least t keys, give $c_i[x]$ an extra key by moving a key from x down into $c_i[x]$, moving a key from $c_i[x]$'s immediate left or right sibling up into x , and moving the appropriate child pointer from the sibling into $c_i[x]$.
- b. If $c_i[x]$ and both of $c_i[x]$'s immediate siblings have $t - 1$ keys, merge $c_i[x]$ with one sibling, which involves moving a key from x down into the new merged node to become the median key for that node.

Deleting Keys: Distribution


Descending down the tree: if k not found in current node x , find the sub-tree $c_i[x]$ that has to contain k .

If $c_i[x]$ has only $t - 1$ keys take action to ensure that we descent to a node of size at least t .


We can encounter two cases.

If $c_i[x]$ has only $t-1$ keys, but a sibling with at least t keys, give $c_i[x]$ an extra key by moving a key from x to $c_i[x]$, moving a key from $c_i[x]$'s immediate left and right sibling up into x , and moving the appropriate child from the sibling into $c_i[x]$ - *distribution*

Deleting Keys: Distribution




Deleting Keys



This is case 3a: ⁹If $c_i[x]$ has only $t - 1$ keys but has an immediate sibling with at least t keys, give $c_i[x]$ an extra key by moving a key from x down into $c_i[x]$, moving a key from $c_i[x]$'s immediate left or right sibling up into x , and moving the appropriate child pointer from the sibling into $c_i[x]$.


Deleting Keys

C is moved to fill B's position and E is moved to fill C's position.




Deleting Keys- Merging

If $c_i[x]$ and both of $c_i[x]$'s siblings have $t - 1$ keys, **merge** c_i with one sibling, which involves moving a key from x down into the new merged node to become the median key for that node



Deleting Keys

Delete D




This is case 3b: If $c_i[x]$ and both of $c_i[x]$'s immediate siblings have $t - 1$ keys, merge $c_i[x]$ with one sibling, which involves moving a key from x down into the new merged node to become the median key for that node.

Deleting Keys

Delete D


the recursion can't descend to node *CL* because it has only 2 keys, so *P* is pushed down and merged with *CL* and *T X* to form *CLPTX*; then, *D* is deleted from a leaf (case 1).



Deleting Keys

Delete D

(e') After (d), the root is deleted and the tree shrinks in height by one.



Deleting Keys

Exercise:

Show the results of deleting *C*, *P*, and *V*, in order, from the tree:

